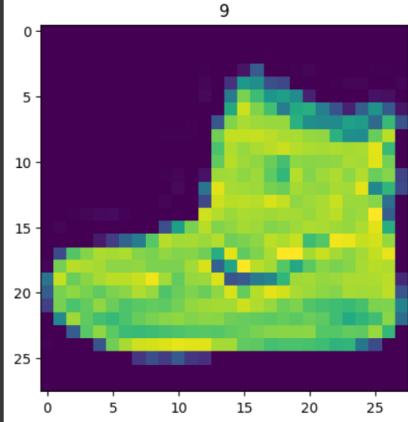


```
0.0000, 0.0000, 0.0235, 0.0000, 0.3862, 0.9569, 0.8706, 0.8627,
0.8549, 0.7961, 0.7765, 0.8667, 0.8431, 0.8353, 0.8706, 0.8627,
0.9608, 0.4667, 0.6549, 0.2196],  
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,  
0.0000, 0.0157, 0.0000, 0.0000, 0.2157, 0.9255, 0.8941, 0.9020,  
0.8941, 0.9412, 0.9098, 0.8353, 0.8549, 0.8745, 0.9176, 0.8510,  
0.8510, 0.8196, 0.3608, 0.0000],  
[0.0000, 0.0000, 0.0039, 0.0157, 0.0235, 0.0275, 0.0078, 0.0000,  
0.0000, 0.0000, 0.0000, 0.0000, 0.9294, 0.8863, 0.8510, 0.8745,
```

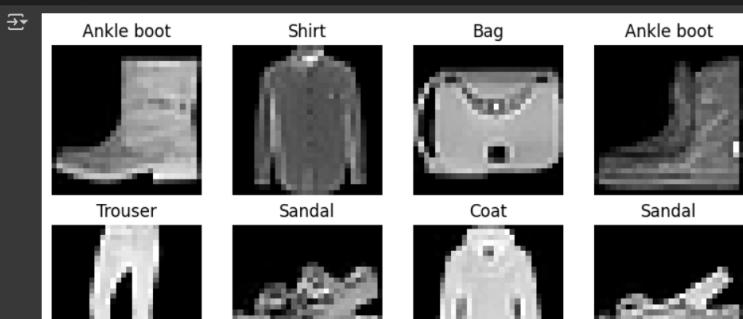
▼ 1.1 Input and output shapes of a computer vision model

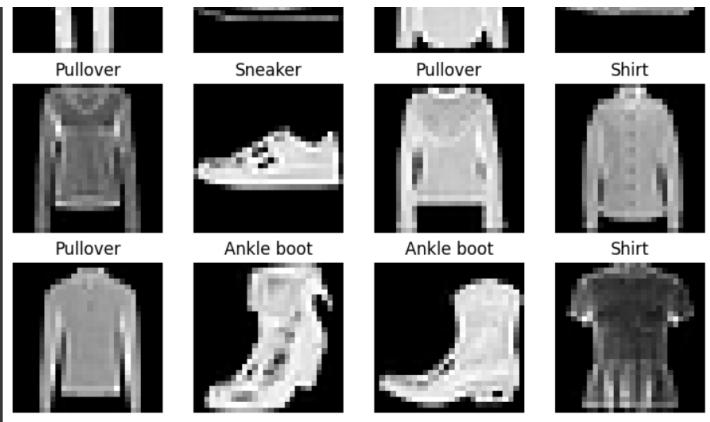
```
✓ 0s [4] 1 image.shape # [color_channels=1, height=28, width=28]  
→ torch.Size([1, 28, 28])  
  
✓ 0s [5] 1 len(train_data.data), len(test_data.data) # we have 60000 training data and 10000 testing data  
→ (60000, 10000)  
  
✓ 0s [6] 1 # checking diff classes  
2  
3 class_names = train_data.classes  
4 class_names  
  
→ ['T-shirt/top',  
'Trouser',  
'Pullover',  
'Dress',  
'Coat',  
'Sandal',  
'Shirt',  
'Sneaker',  
'Bag',  
'Ankle boot']  
  
✓ [7] 1 ## we're working with 10 different classes, it means our problem is multi-class classification.
```

▼ 1.2 Visualizing our data

```
✓ 0s [8] 1 image, label = train_data[0]  
2 print(f"Image Size: {image.shape}")  
3 plt.imshow(image.squeeze())  
4 plt.title(label)  
  
→ Image Size: torch.Size([1, 28, 28])  
Text(0.5, 1.0, '9')  

```

```
✓ 0s [9] 1 # plot more images  
2  
3 torch.manual_seed(42)  
4 fig = plt.figure(figsize=(9, 9))  
5 rows, cols = 4, 4  
6  
7 for i in range (1, rows*cols + 1):  
8     random_idx = torch.randint(0, len(train_data), size=[1]).item()  
9     image, label = train_data[random_idx]  
10    fig.add_subplot(rows, cols, i)  
11    plt.imshow(image.squeeze(), cmap="gray")  
12    plt.title(class_names[label])  
13    plt.axis(False)
```





▼ 2. Prepare DataLoader

It helps load data into a model. For training and for inference.

It turns a large Dataset into a Python iterable of smaller chunks. These smaller chunks are called batches or mini-batches and can be set by the batch_size parameter.

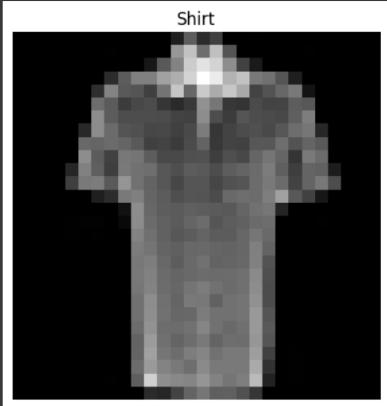
```
[10] 1 from torch.utils.data import DataLoader
2
3 # set the batch_size hyper parameter
4 BATCH_SIZE = 32
5
6 # turn datasets into iterables
7 train_dataloader = DataLoader(train_data, # data you want to turn into iterables
8                               batch_size = BATCH_SIZE,
9                               shuffle = True) # shuffle data every epoch
10
11 test_dataloader = DataLoader(test_data,
12                               batch_size = BATCH_SIZE,
13                               shuffle = False) # we dont necessarily have to shuffle the testing data
```

```
[11] 1 len(train_dataloader), len(test_dataloader)
2
3 ➜ (1875, 313)
```

```
[12] 1 # Check out what's inside the training dataloader
2 train_features_batch, train_labels_batch = next(iter(train_dataloader))
3 train_features_batch.shape, train_labels_batch.shape
4
5 ➜ (torch.Size([32, 1, 28, 28]), torch.Size([32]))
```

```
[13] 1 ## And we can see that the data remains unchanged by checking a single sample.
2
3 # Show a sample
4 torch.manual_seed(42)
5
6 random_idx = torch.randint(0, len(train_features_batch), size=[1]).item()
7 img, label = train_features_batch[random_idx], train_labels_batch[random_idx]
8 plt.imshow(img.squeeze(), cmap="gray")
9 plt.title(class_names[label])
10 plt.axis("off");
11 print(f"Image size: {img.shape}")
12 print(f"Label: {label}, label size: {label.shape}")
```

```
6 ➜ Image size: torch.Size([1, 28, 28])
7 Label: 6, label size: torch.Size([])
```



▼ 3. Model 0: Build a baseline model

```
[14] 1 # create a flatten layer
2 flatten_model = nn.Flatten() # nn.Flatten() compresses the dimensions of a tensor into a single vector
3
4 # get a single sample
5 x = train_features_batch[0]
```

```

6
7 # flatten the sample
8 output = flatten_model(x)
9
10 # lets see now
11 print(f"Shape before flattening: {x.shape}")
12 print(f"Shape after flattening: {output.shape}")
13

```

⇒ Shape before flattening: torch.Size([1, 28, 28])
Shape after flattening: torch.Size([1, 784])

▼ Why do this?

Why do this?

Because we've now turned our pixel data from height and width dimensions into one long feature vector.

And nn.Linear() layers like their inputs to be in the form of feature vectors.

```

✓ [15] 1 from torch import nn
2 class FashionMNISTModelV0(nn.Module):
3     def __init__(self, input_shape: int, hidden_units: int, output_shape: int): # takes 3 args
4         super().__init__()
5         self.layer_stack = nn.Sequential(
6             nn.Flatten(), # neural networks like their inputs in vector form
7             nn.Linear(in_features=input_shape, out_features=hidden_units), # in_features = number of features in a data sample (784 pixels)
8             nn.Linear(in_features=hidden_units, out_features=output_shape)
9         )
10
11     def forward(self, x):
12         return self.layer_stack(x)
13
14 torch.manual_seed(42)
15
16 # Need to setup model with input parameters
17 model_0 = FashionMNISTModelV0(input_shape=784, # one for every pixel (28x28)
18                                hidden_units=10, # how many units in the hidden layer
19                                output_shape=len(class_names) # one for every class
20 )
21 model_0.to("cpu") # keep model on CPU to begin with

```

⇒ FashionMNISTModelV0(
(layer_stack): Sequential(
 (0): Flatten(start_dim=1, end_dim=-1)
 (1): Linear(in_features=784, out_features=10, bias=True)
 (2): Linear(in_features=10, out_features=10, bias=True)
)

▼ 3.1 Setup loss, optimizer and evaluation metrics

```

✓ [16] 1 import requests
2 from pathlib import Path
3
4 # Download helper functions from Learn PyTorch repo (if not already downloaded)
5 if Path("helper_functions.py").is_file():
6     print("helper_functions.py already exists, skipping download")
7 else:
8     print("downloading helper_functions.py")
9     # Note: you need the "raw" GitHub URL for this to work
10    request = requests.get("https://raw.githubusercontent.com/mrdbourke/pytorch-deep-learning/main/helper_functions.py")
11    with open("helper_functions.py", "wb") as f:
12        f.write(request.content)

```

⇒ Downloading helper_functions.py

```

✓ [17] 1 # Import accuracy metric
2 from helper_functions import accuracy_fn # Note: could also use torchmetrics.Accuracy(task = 'multiclass', num_classes=len(class_names)).to(device)
3
4 # Setup loss function and optimizer
5 loss_fn = nn.CrossEntropyLoss() # this is also called "criterion"/"cost function" in some places
6 optimizer = torch.optim.SGD(params=model_0.parameters(), lr=0.1)

```

▼ 3.2 Creating a function to time our experiments

```

✓ [18] 1 from timeit import default_timer as timer
2 def print_train_time(start: float, end: float, device: torch.device = None):
3     """Prints difference between start and end time.
4
5     Args:
6         start (float): Start time of computation (preferred in timeit format).
7         end (float): End time of computation.
8         device ([type], optional): Device that compute is running on. Defaults to None.
9
10    Returns:
11        float: time between start and end in seconds (higher is longer).
12    """
13    total_time = end - start
14    print(f"Train time on {device}: {total_time:.3f} seconds")
15    return total_time

```

▼ 3.3 Creating a training loop and training a model on batches of data

Let's step through it:

- Loop through epochs.
- Loop through training batches, perform training steps, calculate the train_loss per batch.

Loop through training batches, perform training steps, calculate the train loss per batch.

- Loop through testing batches, perform testing steps, calculate the test loss per batch.
- Print out what's happening.
- Time it all (for fun).

```
[19]  1 # Import tqdm for progress bar
2 from tqdm.auto import tqdm
3
4 # Set the seed and start the timer
5 torch.manual_seed(42)
6 train_time_start_on_cpu = timer()
7
8 # Set the number of epochs (we'll keep this small for faster training times)
9 epochs = 3
10
11 # Create training and testing loop
12 for epoch in tqdm(range(epochs)):
13     print(f'Epoch: {epoch}\n-----')
14     ### Training
15     train_loss = 0
16     # Add a loop to loop through training batches
17     for batch, (X, y) in enumerate(train_dataloader):
18         model_0.train()
19         # 1. Forward pass
20         y_pred = model_0(X)
21
22         # 2. Calculate loss (per batch)
23         loss = loss_fn(y_pred, y)
24         train_loss += loss # accumulatively add up the loss per epoch
25
26         # 3. Optimizer zero grad
27         optimizer.zero_grad()
28
29         # 4. Loss backward
30         loss.backward()
31
32         # 5. Optimizer step
33         optimizer.step()
34
35         # Print out how many samples have been seen
36         if batch % 400 == 0:
37             print(f"Looked at {batch * len(X)}/{len(train_dataloader.dataset)} samples")
38
39     # Divide total train loss by length of train dataloader (average loss per batch per epoch)
40     train_loss /= len(train_dataloader)
41
42     ### Testing
43     # Setup variables for accumulatively adding up loss and accuracy
44     test_loss, test_acc = 0, 0
45     model_0.eval()
46     with torch.inference_mode():
47         for X, y in test_dataloader:
48             # 1. Forward pass
49             test_pred = model_0(X)
50
51             # 2. Calculate loss (accumulatively)
52             test_loss += loss_fn(test_pred, y) # accumulatively add up the loss per epoch
53
54             # 3. Calculate accuracy (preds need to be same as y_true)
55             test_acc += accuracy_fn(y_true=y, y_pred=test_pred.argmax(dim=1))
56
57             # Calculations on test metrics need to happen inside torch.inference_mode()
58             # Divide total test loss by length of test dataloader (per batch)
59             test_loss /= len(test_dataloader)
60
61             # Divide total accuracy by length of test dataloader (per batch)
62             test_acc /= len(test_dataloader)
63
64     ## Print out what's happening
65     print(f'\nTrain loss: {train_loss:.5f} | Test loss: {test_loss:.5f}, Test acc: {test_acc:.2f}%\n')
66
67 # Calculate training time
68 train_time_end_on_cpu = timer()
69 total_train_time_model_0 = print_train_time(start=train_time_start_on_cpu,
70                                              end=train_time_end_on_cpu,
71                                              device=str(next(model_0.parameters()).device))
```

⌚ 100% [3/3] 00:28<00:00, 9.49s/it]

Epoch: 0

Looked at 0/60000 samples
Looked at 12800/60000 samples
Looked at 25600/60000 samples
Looked at 38400/60000 samples
Looked at 51200/60000 samples

Train loss: 0.59039 | Test loss: 0.50954, Test acc: 82.04%

Epoch: 1

Looked at 0/60000 samples
Looked at 12800/60000 samples
Looked at 25600/60000 samples
Looked at 38400/60000 samples
Looked at 51200/60000 samples

Train loss: 0.47633 | Test loss: 0.47989, Test acc: 83.20%

Epoch: 2

Looked at 0/60000 samples
Looked at 12800/60000 samples
Looked at 25600/60000 samples
Looked at 38400/60000 samples
Looked at 51200/60000 samples

Train loss: 0.45503 | Test loss: 0.47664, Test acc: 83.43%

Train time on cpu: 28.527 seconds

▼ 4. Make predictions and get Model 0 results

```
[20] 1 torch.manual_seed(42)
2 def eval_model(model: torch.nn.Module,
3                 data_loader: torch.utils.data.DataLoader,
4                 loss_fn: torch.nn.Module,
5                 accuracy_fn):
6     """Returns a dictionary containing the results of model predicting on data_loader.
7
8     Args:
9         model (torch.nn.Module): A Pytorch model capable of making predictions on data_loader.
10        data_loader (torch.utils.data.DataLoader): The target dataset to predict on.
11        loss_fn (torch.nn.Module): The loss function of model.
12        accuracy_fn: An accuracy function to compare the models predictions to the truth labels.
13
14    Returns:
15        (dict): Results of model making predictions on data_loader.
16    """
17    loss, acc = 0, 0
18    model.eval()
19    with torch.inference_mode():
20        for X, y in data_loader:
21            # Make predictions with the model
22            y_pred = model(X)
23
24            # Accumulate the loss and accuracy values per batch
25            loss += loss_fn(y_pred, y)
26            acc += accuracy_fn(y_true=y,
27                                y_pred=y_pred.argmax(dim=1)) # For accuracy, need the prediction labels (logits -> pred_prob -> pred_labels)
28
29    # Scale loss and acc to find the average loss/acc per batch
30    loss /= len(data_loader)
31    acc /= len(data_loader)
32
33    return {"model_name": model.__class__.__name__, # only works when model was created with a class
34            "model_loss": loss.item(),
35            "model_acc": acc}
36
37 # Calculate model 0 results on test dataset
38 model_0_results = eval_model(model=model_0, data_loader=test_dataloader,
39                             loss_fn=loss_fn, accuracy_fn=accuracy_fn
40 )
41 model_0_results
```

⇒ {'model_name': 'FashionMNISTModelV0',
'model_loss': 0.47663894295692444,
'model_acc': 83.42651757188499}

▼ 5. Setup device agnostic-code (for using a GPU if there is one)

```
[21] 1 # Setup device agnostic code
2 import torch
3 device = "cuda" if torch.cuda.is_available() else "cpu"
4 device
```

⇒ 'cpu'

▼ 6. Model 1: Building a better model with non-linearity

```
[22] 1 # Create a model with non-linear and linear layers
2 class FashionMNISTModelV1(nn.Module):
3     def __init__(self, input_shape: int, hidden_units: int, output_shape: int):
4         super().__init__()
5         self.layer_stack = nn.Sequential(
6             nn.Flatten(), # flatten inputs into single vector
7             nn.Linear(in_features=input_shape, out_features=hidden_units),
8             nn.ReLU(),
9             nn.Linear(in_features=hidden_units, out_features=output_shape),
10            nn.ReLU()
11        )
12
13    def forward(self, x: torch.Tensor):
14        return self.layer_stack(x)
```

```
[23] 1 torch.manual_seed(42)
2 model_1 = FashionMNISTModelV1(input_shape=784, # number of input features
3                               hidden_units=10,
4                               output_shape=len(class_names) # number of output classes desired
5 ).to(device) # send model to GPU if it's available
6 next(model_1.parameters()).device # check model device
```

⇒ device(type='cpu')

▼ 6.1 Setup loss, optimizer and evaluation metrics

```
[24] 1 from helper functions import accuracy_fn
2 loss_fn = nn.CrossEntropyLoss()
3 optimizer = torch.optim.SGD(params=model_1.parameters(),
4                             lr=0.1)
```

▼ 6.2 Functionizing training and test loops

```

✓ [25] 1 def train_step(model: torch.nn.Module,
2                 data_loader: torch.utils.data.DataLoader,
3                 loss_fn: torch.nn.Module,
4                 optimizer: torch.optim.Optimizer,
5                 accuracy_fn,
6                 device: torch.device = device):
7     train_loss, train_acc = 0, 0
8     model.to(device)
9     for batch, (X, y) in enumerate(data_loader):
10        # Send data to GPU
11        X, y = X.to(device), y.to(device)
12
13        # 1. Forward pass
14        y_pred = model(X)
15
16        # 2. Calculate loss
17        loss = loss_fn(y_pred, y)
18        train_loss += loss
19        train_acc += accuracy_fn(y_true=y,
20                                y_pred=y_pred.argmax(dim=1)) # Go from logits -> pred labels
21
22        # 3. Optimizer zero grad
23        optimizer.zero_grad()
24
25        # 4. Loss backward
26        loss.backward()
27
28        # 5. Optimizer step
29        optimizer.step()
30
31    # Calculate loss and accuracy per epoch and print out what's happening
32    train_loss /= len(data_loader)
33    train_acc /= len(data_loader)
34    print(f"Train loss: {train_loss:.5f} | Train accuracy: {train_acc:.2f}%")
35
36 def test_step(data_loader: torch.utils.data.DataLoader,
37               model: torch.nn.Module,
38               loss_fn: torch.nn.Module,
39               accuracy_fn,
40               device: torch.device = device):
41     test_loss, test_acc = 0, 0
42     model.to(device)
43     model.eval() # put model in eval mode
44     # Turn on inference context manager
45     with torch.inference_mode():
46       for X, y in data_loader:
47         # Send data to GPU
48         X, y = X.to(device), y.to(device)
49
50         # 1. Forward pass
51         test_pred = model(X)
52
53         # 2. Calculate loss and accuracy
54         test_loss += loss_fn(test_pred, y)
55         test_acc += accuracy_fn(y_true=y,
56                                y_pred=test_pred.argmax(dim=1)) # Go from logits -> pred labels
57
58     # Adjust metrics and print out
59     test_loss /= len(data_loader)
60     test_acc /= len(data_loader)
61     print(f"Test loss: {test_loss:.5f} | Test accuracy: {test_acc:.2f}%\n")

```

```

✓ [26] 1 torch.manual_seed(42)
2
3 # Measure time
4 from timeit import default_timer as timer
5 train_time_start_on_gpu = timer()
6
7 epochs = 3
8 for epoch in tqdm(range(epochs)):
9   print(f"Epoch: {epoch}\n-----")
10  train_step(data_loader=train_dataloader,
11             model=model_1,
12             loss_fn=loss_fn,
13             optimizer=optimizer,
14             accuracy_fn=accuracy_fn
15           )
16  test_step(data_loader=test_dataloader,
17            model=model_1,
18            loss_fn=loss_fn,
19            accuracy_fn=accuracy_fn
20          )
21
22 train_time_end_on_gpu = timer()
23 total_train_time_model_1 = print_train_time(start=train_time_start_on_gpu,
24                                              end=train_time_end_on_gpu,
25                                              device=device)

```

→ 100%  3/3 [00:28<00:00, 9.37s/it]

Epoch: 0

Train loss: 1.09199 | Train accuracy: 61.34%

Test loss: 0.95636 | Test accuracy: 65.00%

Epoch: 1

Train loss: 0.78101 | Train accuracy: 71.93%

Test loss: 0.72227 | Test accuracy: 73.91%

Epoch: 2

Train loss: 0.67027 | Train accuracy: 75.94%

Test loss: 0.68500 | Test accuracy: 75.02%

Train time on cpu: 28.042 seconds

```
[32] 1 torch.manual_seed(42)
2
3 # Note: This will error due to `eval_model()` not using device agnostic code
4 model_1_results = eval_model(model=model_1,
5 | data_loader=test_dataloader,
6 | loss_fn=loss_fn,
7 | accuracy_fn=accuracy_fn)
8 model_1_results

→ {'model_name': 'FashionMNISTModelV1',
'model_loss': 0.6850009560585022,
'model_acc': 75.01996805111821}
```

▼ 7. Model 2: Building a Convolutional Neural Network (CNN)

```
✓ [27] 1 # Create a convolutional neural network
2 class FashionMNISTModelV2(nn.Module):
3     """
4     Model architecture copying TinyVGG from:
5     https://poloclub.github.io/cnn-explainer/
6     """
7     def __init__(self, input_shape: int, hidden_units: int, output_shape: int):
8         super().__init__()
9         self.block_1 = nn.Sequential(
10             nn.Conv2d(in_channels=input_shape,
11                       out_channels=hidden_units,
12                       kernel_size=3, # how big is the square that's going over the image?
13                       stride=1, # default
14                       padding=1),# options = "valid" (no padding) or "same" (output has same shape as input) or int for specific number
15             nn.ReLU(),
16             nn.Conv2d(in_channels=hidden_units,
17                       out_channels=hidden_units,
18                       kernel_size=3,
19                       stride=1,
20                       padding=1),
21             nn.ReLU(),
22             nn.MaxPool2d(kernel_size=2,
23                         | | | | | stride=2) # default stride value is same as kernel_size
24         )
25         self.block_2 = nn.Sequential(
26             nn.Conv2d(hidden_units, hidden_units, 3, padding=1),
27             nn.ReLU(),
28             nn.Conv2d(hidden_units, hidden_units, 3, padding=1),
29             nn.ReLU(),
30             nn.MaxPool2d(2)
31         )
32         self.classifier = nn.Sequential(
33             nn.Flatten(),
34             # Where did this in_features shape come from?
35             # It's because each layer of our network compresses and changes the shape of our input data.
36             nn.Linear(in_features=hidden_units*7*7,
37                         | | | | | out_features=output_shape)
38         )
39
40     def forward(self, x: torch.Tensor):
41         x = self.block_1(x)
42         # print(x.shape)
43         x = self.block_2(x)
44         # print(x.shape)
45         x = self.classifier(x)
46         # print(x.shape)
47         return x
48
49 torch.manual_seed(42)
50 model_2 = FashionMNISTModelV2(input_shape=1,
51 | hidden_units=10,
52 | output_shape=len(class_names)).to(device)
53 model_2
```

```
→ FashionMNISTModelV2(
    (block_1): Sequential(
        (0): Conv2d(1, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
        (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (3): ReLU()
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (block_2): Sequential(
        (0): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
        (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (3): ReLU()
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Flatten(start_dim=1, end_dim=-1)
        (1): Linear(in_features=490, out_features=10, bias=True)
    )
)
```

▼ 7.1 Setup a loss function and optimizer for model_2

```
✓ [28] 1 # Setup loss and optimizer
2 loss_fn = nn.CrossEntropyLoss()
3 optimizer = torch.optim.SGD(params=model_2.parameters(),
4 | | | | | lr=0.1)
```

▼ 7.2 Training and testing model_2 using our training and test functions

```

✓ [29] 1 torch.manual_seed(42)
2
3 # Measure time
4 from timeit import default_timer as timer
5 train_time_start_model_2 = timer()
6
7 # Train and test model
8 epochs = 3
9 for epoch in tqdm(range(epochs)):
10    print(f"Epoch: {epoch}\n-----")
11    train_step(data_loader=train_dataloader,
12               model=model_2,
13               loss_fn=loss_fn,
14               optimizer=optimizer,
15               accuracy_fn=accuracy_fn,
16               device=device
17    )
18    test_step(data_loader=test_dataloader,
19               model=model_2,
20               loss_fn=loss_fn,
21               accuracy_fn=accuracy_fn,
22               device=device
23    )
24
25 train_time_end_model_2 = timer()
26 total_train_time_model_2 = print_train_time(start=train_time_start_model_2,
27                                              end=train_time_end_model_2,
28                                              device=device)

```

→ 100% 3/3 [02:19<00:00, 46.34s/i]

Epoch: 0

Train loss: 0.58653 | Train accuracy: 78.74%

Test loss: 0.39255 | Test accuracy: 86.04%

Epoch: 1

Train loss: 0.36231 | Train accuracy: 86.89%

Test loss: 0.35723 | Test accuracy: 86.70%

Epoch: 2

Train loss: 0.32485 | Train accuracy: 88.22%

Test loss: 0.32124 | Test accuracy: 88.02%

Train time on cpu: 139.285 seconds

```

✓ [30] 1 # Get model_2 results
2 model_2_results = eval_model(
3     model=model_2,
4     data_loader=test_dataloader,
5     loss_fn=loss_fn,
6     accuracy_fn=accuracy_fn
7 )
8 model_2_results

```

→ {'model_name': 'FashionMNISTModelV2',
 'model_loss': 0.3212365508079529,
 'model_acc': 88.01916932907348}

8. Compare model results and training time

```

✓ [33] 1 import pandas as pd
2 compare_results = pd.DataFrame([model_0_results, model_1_results, model_2_results])
3 compare_results

```

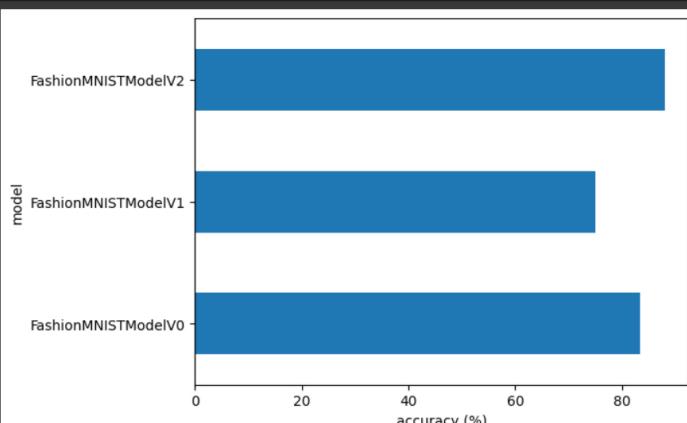
→

	model_name	model_loss	model_acc
0	FashionMNISTModelV0	0.476639	83.426518
1	FashionMNISTModelV1	0.685001	75.019968
2	FashionMNISTModelV2	0.321237	88.019169

```

✓ [34] 1 # Visualize our model results
2 compare_results.set_index("model_name")["model_acc"].plot(kind="barh")
3 plt.xlabel("accuracy (%)")
4 plt.ylabel("model");

```



- ▼ 9. Make and evaluate random predictions with best model

```
[35] 1 def make_predictions(model: torch.nn.Module, data: list, device: torch.device = device):
2     pred_probs = []
3     model.eval()
4     with torch.inference_mode():
5         for sample in data:
6             # Prepare sample
7             sample = torch.unsqueeze(sample, dim=0).to(device) # Add an extra dimension and send sample to device
8
9             # Forward pass (model outputs raw logit)
10            pred_logit = model(sample)
11
12            # Get prediction probability (logit -> prediction probability)
13            pred_prob = torch.softmax(pred_logit.squeeze(), dim=0) # note: perform softmax on the "logits" dimension, not "batch" dimension (in this case we have a batch size of 1,
14
15            # Get pred_prob off GPU for further calculations
16            pred_probs.append(pred_prob.cpu())
17
18    # Stack the pred_probs to turn list into a tensor
19    return torch.stack(pred_probs)
```

```
[36] 1 import random
2 random.seed(42)
3 test_samples = []
4 test_labels = []
5 for sample, label in random.sample(list(test_data), k=9):
6     | test_samples.append(sample)
7     | test_labels.append(label)
8
9 # View the first test sample shape and label
10 print(f"Test sample image shape: {test_samples[0].shape}\nTest sample label: {test_labels[0]} ({class_names[test_labels[0]]})")
```

```
tensor([[1.7889e-07, 3.0100e-07, 7.4711e-08, 2.9112e-07, 4.0070e-08, 9.9988e-01,  
        9.7217e-07, 1.5529e-05, 4.1133e-05, 6.1028e-05],  
       [8.4885e-02, 5.9999e-01, 1.6154e-03, 1.5859e-01, 3.7294e-02, 1.3594e-04,  
        1.2469e-01, 3.1763e-04, 1.2822e-03, 1.8921e-04]])
```

```
tensor([[1.7889e-07, 3.0100e-07, 7.4711e-08, 2.9112e-07, 4.0070e-08, 9.9988e-01,  
        9.7217e-07, 1.5529e-05, 4.1133e-05, 6.1028e-05],  
       [8.4885e-02, 5.9999e-01, 1.6154e-03, 1.5859e-01, 3.7294e-02, 1.3594e-04,  
        1.2469e-01, 3.1763e-04, 1.2822e-03, 1.8921e-04]])
```

```
✓ [39] 1 # Turn the prediction probabilities into prediction labels by taking the argmax()
  2 pred_classes = pred_probs.argmax(dim=1)
  3 pred_classes
```

→ tensor([5, 1, 7, 4, 3, 0, 4, 7, 1])

[40] 1 # Are our predictions in the same form

```
2 test_labels, pred_classes
```

```
2 plt.figure(figsize=(9, 9))
3 nrows = 3
4 ncols = 3
5 for i, sample in enumerate(test_samples):
6     # Create a subplot
7     plt.subplot(nrows, ncols, i+1)
8
9     # Plot the target image
10    plt.imshow(sample.squeeze(), cmap="gray")
11
12    # Find the prediction label (in text form, e.g. "Sandal")
13    pred_label = class_names[pred_classes[i]]
14
15    # Get the truth label (in text form, e.g. "T-shirt")
16    truth_label = class_names[test_labels[i]]
17
18    # Create the title text of the plot
19    title_text = f"Pred: {pred_label} | Truth: {truth_label}"
20
21    # Check for equality and change title colour accordingly
22    if pred_label == truth_label:
23        plt.title(title_text, fontsize=10, c="g") # green text if correct
24    else:
25        plt.title(title_text, fontsize=10, c="r") # red text if wrong
26
27    plt.axis(False);
```



▼ 10. Making a confusion matrix for further prediction evaluation

```
[42] 1 # Import tqdm for progress bar
2 from tqdm.auto import tqdm
3
4 # 1. Make predictions with trained model
5 y_preds = []
6 model_2.eval()
7 with torch.inference_mode():
8     for X, y in tqdm(test_dataloader, desc="Making predictions"):
9         # Send data and targets to target device
10        X, y = X.to(device), y.to(device)
11        # Do the forward pass
12        y_logit = model_2(X)
13        # Turn predictions from logits -> prediction probabilities -> predictions labels
14        y_pred = torch.softmax(y_logit, dim=1).argmax(dim=1) # note: perform softmax on the "logits" dimension, not "batch" dimension (in this case we have a batch size of 32, so can just use dim=1)
15        # Put predictions on CPU for evaluation
16        y_preds.append(y_pred.cpu())
17 # Concatenate list of predictions into a tensor
18 y_pred_tensor = torch.cat(y_preds)
```

⌚ Making predictions: 100% [██████████] 313/313 [00:03<00:00, 90.53it/s]

```
[43] 1 # See if torchmetrics exists, if not, install it
2 try:
3     import torchmetrics, mlxtend
4     print(f"mlxtend version: {mlxtend.__version__}")
5     assert int(mlxtend.__version__.split(".")[1]) >= 19, "mlxtend version should be 0.19.0 or higher"
6 except:
7     !pip install -q torchmetrics -U mlxtend # <- Note: If you're using Google Colab, this may require restarting the runtime
8     import torchmetrics, mlxtend
9     print(f"mlxtend version: {mlxtend.__version__}")

mlxtend version: 0.23.4
```

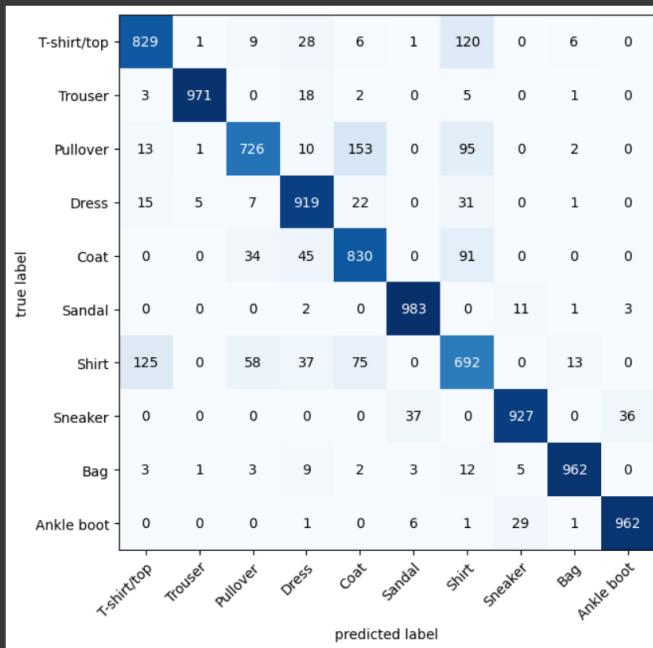
⌚ 0.23.4

```
[44] 1 # Import mlxtend upgraded version
2 import mlxtend
3 print(mlxtend.__version__)
4 assert int(mlxtend.__version__.split(".")[1]) >= 19 # should be version 0.19.0 or higher
```

⌚ 0.23.4

```
[45] 1 from torchmetrics import confusionMatrix
2 from mlxtend.plotting import plot_confusion_matrix
3
4 # 2. Setup confusion matrix instance and compare predictions to targets
5 confmat = ConfusionMatrix(num_classes=len(class_names), task='multiclass')
6 confmat_tensor = confmat(preds=y_pred_tensor,
7                         target=test_data.targets)
8
9 # 3. Plot the confusion matrix
10 fig, ax = plot_confusion_matrix(
11     conf_mat=confmat_tensor.numpy(), # matplotlib likes working with NumPy
12     class_names=class_names, # include the row and column labels into class names
```

```
12 |     class_names=class_names, # turn the row and column labels into class names
13 |     figsize=(10, 7)
14 | );
```



```
[51] 1 !pip install nbstripout
2
```

```
→ Requirement already satisfied: nbstripout in /usr/local/lib/python3.11/dist-packages (0.8.1)
Requirement already satisfied: nbformat in /usr/local/lib/python3.11/dist-packages (from nbstripout) (5.10.4)
Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib/python3.11/dist-packages (from nbformat->nbstripout) (2.21.1)
Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/python3.11/dist-packages (from nbformat->nbstripout) (4.24.0)
Requirement already satisfied: jupyter-core!=5.0.*,>4.12 in /usr/local/lib/python3.11/dist-packages (from nbformat->nbstripout) (5.8.1)
Requirement already satisfied: traitlets>=5.1 in /usr/local/lib/python3.11/dist-packages (from nbformat->nbstripout) (5.7.1)
Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.11/dist-packages (from jsonschema>=2.6->nbformat->nbstripout) (25.3.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /usr/local/lib/python3.11/dist-packages (from jsonschema>=2.6->nbformat->nbstripout) (2025.4.1)
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/python3.11/dist-packages (from jsonschema>=2.6->nbformat->nbstripout) (0.36.2)
Requirement already satisfied: rpyd-py>=0.7.1 in /usr/local/lib/python3.11/dist-packages (from jsonschema>=2.6->nbformat->nbstripout) (0.25.1)
Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/python3.11/dist-packages (from jupyter-core!=5.0.*,>4.12->nbformat->nbstripout) (4.3.8)
Requirement already satisfied: typing-extensions>=4.4.0 in /usr/local/lib/python3.11/dist-packages (from referencing>=0.28.4->jsonschema>=2.6->nbformat->nbstripout) (4.14.0)
```

```
[52] 1 !nbstripout your_notebook.ipynb
2
```

```
→ Could not strip 'your_notebook.ipynb': file not found
```

```
[53] 1 !nbstripout CNN_from_scratch.ipynb
```

```
→ Could not strip 'CNN_from_scratch.ipynb': file not found
```

```
[54] 1 ll
```

```
→ data helper_functions.py __pycache__ sample_data
```

```
[55] 1 from pathlib import Path
2
3 # Create models directory (if it doesn't already exist), see: https://docs.python.org/3/library/pathlib.html#pathlib.Path.mkdir
4 MODEL_PATH = Path("models")
5 MODEL_PATH.mkdir(parents=True, # create parent directories if needed
6 | | | | | exist_ok=True # if models directory already exists, don't error
7 )
8
9 # Create model save path
10 MODEL_NAME = "03_pytorch_computer_vision_model_2.pth"
11 MODEL_SAVE_PATH = MODEL_PATH / MODEL_NAME
12
13 # Save the model state dict
14 print(f"Saving model to: {MODEL_SAVE_PATH}")
15 torch.save(model_2.state_dict(), # only saving the state_dict() only saves the learned parameters
16 | | | | f=MODEL_SAVE_PATH)
```

```
→ Saving model to: models/03_pytorch_computer_vision_model_2.pth
```

```
[72] 1 Start coding or generate with AI.
```

```
→ Could not strip 'content/models/03_pytorch_computer_vision_model.ipynb': file not found
```

```
[ ] 1 Start coding or generate with AI.
```

[Colab paid products](#) - [Cancel contracts here](#)

 Variables  Terminal



✓ 8:48 AM  Python 3